

# In this lecture, we will discuss...

- ✧ Rationale behind NoSQL
- ✧ Scaling Issues in RDBMS
- ✧ NoSQL: What is it?




# Why RDBMS

- ✧ Relational Databases – popular and commonly used
- ✧ Initially designed for non distributed
- ✧ Low Cost RDBMS alternatives (PostgreSQL, MySQL, SQLite )
- ✧ Very Transactional - across tables and commands, and can even be transactional across distributed resources (XA) -- at a cost
- ✧ Supports Joins -- across multiple tables allowing for normalized forms of data to be stored once



# Why NoSQL

- ✧ Explosion in data
- ✧ Object/Relational Impedance mismatch 
  - Objects are **constantly being moved** in/out of tables/rows
- ✧ RDBMS normalization and joins are **powerful**, but add up in **cost**
  - Complex objects stored across many tables and rows can be **expensive** to handle


# Why NoSQL

- ✧ “Big” data handling with better performance
- ✧ Supports unstructured data
  - Unique data type extensions can be easily integrated into existing collections
- ✧ Operational issues (scale, performance and availability)




# Scaling Out

## Vertical Scaling



## Horizontal Scaling



# What is NoSQL

- ✧ Stands for “Not Only SQL”
- ✧ No Fixed Schema
- ✧ Non-relational data storage systems



# Summary

- ✧ NoSQL – very popular and major companies especially social networking sites such as Twitter, Facebook, LinkedIn, and Digg use NoSQL DB
- ✧ Excellent performance and stability, fast and scalable and fairly simple model
- ✧ Supports unstructured format, which makes it very agile
- ✧ NoSQL is mostly gained when access patterns to complex objects are understood and modeled correctly up front



# What's Next?

## Categories of NoSQL





# In this lecture, we will discuss...

- ✧ Categories of NoSQL
- ✧ NoSQL vs. RDBMS

# Categories of NoSQL – Key/Value

✧ Value can be String or JSON

✧ Key-value hash

✧ Solutions 

- Dynamo
- Redis
- Memcached

ID	Attributes
1234	John Doe
1235	{ "Name": "Godfather", "Genre": "Drama", "Actor": "Robert DeNiro", "Director": "Francis Ford Coppola" }

# Categories of NoSQL – Document

- ✧ Stores documents based up of tagged elements
- ✧ Persistent and query-able
- ✧ Solutions
  - MongoDB
  - CouchDB

```
{  
  "id": 1234,  
  "name": "Departed",  
  "actors": [  
    {  
      "actor": "Leonardo DeCaprio"  
    },  
    {  
      "actor": "Jack Nicholson"  
    }  
  ],  
  "director": "Martin Scorsese",  
  "genre": "drama"  
}
```

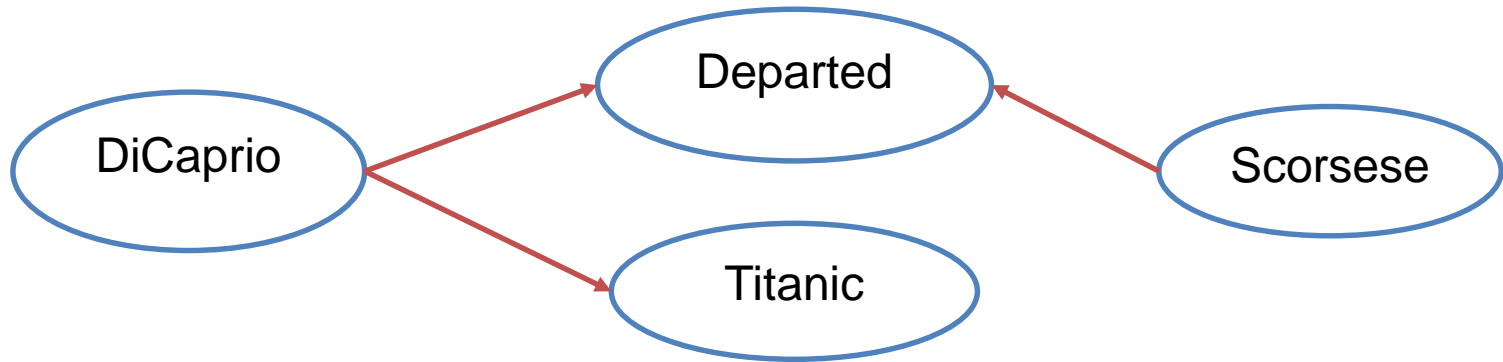
# Categories of NoSQL – Column

- ✧ Uses flat structure, but with keys stored in **columns** rather than rows:
- ✧ Solutions
  - Cassandra
  - Hbase

ID	101	102	103
Name	The Godfather	The Departed	Titanic
Actor	Leonardo DiCaprio	Al Pacino	Leonardo DiCaprio
Director	Francis Ford Coppola	Martin Scorsese	James Cameron

# Categories of NoSQL – Graph

- ✧ A network database that uses **edges and nodes** to **represent and store** data
- ✧ Solutions
  - Neo4J



## NoSQL – Not supported

✧ Joins are not supported

- Embedded documents or in middle tier code 

## ✦💬 ACID Transactions

- Supported at a document level only

```

1 {
2     "title": "The Departed",
3     "type": "Movie",
4     "director": "Martin Scorsese",
5     "actors": [
6         {
7             "actorName": "Leonardo DiCaprio",
8             "character": "Billy",
9             "main": true,
10            "urlCharacter": "http://www.imdb.com/character/ch0251381",
11            "urlPhoto": "http://ia.media-imdb.com/images/M/MV5BMjI0MTg3MzI0M158M158anBnXkFtZ...
12            "urlProfile": "http://www.imdb.com/name/nm0000138"
13        },
14        {
15            "actorName": "Matt Damon",
16            "character": "Colin Sullivan",
17            "main": true,
18            "urlCharacter": "http://www.imdb.com/character/ch0002488",
19            "urlPhoto": "http://ia.media-imdb.com/images/M/MV5BMjM0NzYzNDgxM158M158anBnXkFtZ...
20            "urlProfile": "http://www.imdb.com/name/nm0000354"
21        }
22    ]
23 }

```



# NoSQL vs RDBMS – How to pick?

## ✧ Nature of data

- Row/column (structured) – RDBMS
- Unstructured, complex (geo-spatial or engineering data) which needs nesting - NoSQL

## ✧ Schema

- Static – RDBMS, Dynamic – NoSQL

# NoSQL vs RDBMS – How to pick?

- ✧ Self contained – NoSQL, Joins – RDBMS
- ✧ Flexibility of query
  - RDBMS Joins allow for flexibility
  - NoSQL - Duplication of data, implement joins in middle-ware





# Summary

- ✧ 4 different categories offering different choices
- ✧ Pick what is best for your application (Relational or NoSQL)



# What's Next?

MongoDB



# In this lecture, we will discuss...

- ✧ What MongoDB is
- ✧ Reasons to use MongoDB

# What is MongoDB

- ✧ Created by 10gen (term coined from humongous)
- ✧ Definition:
  - MongoDB is an open source, document-oriented database designed with both scalability and developer agility in mind
- ✧ Storage: JSON-like documents and “schemaless”
- ✧ Well suited for Object Oriented programming



# What is MongoDB?

- ✧ Stores data in **BSON** format (Binary JSON)
- ✧ Binary form for representing **simple data structures** and **associative arrays**

```
1 {
2   "_id": 101,
3   "title": "The Departed",
4   "type": "Movie",
5   "director": "Martin Scorsese",
6   "actors": [
7     {
8       "actorName": "Leonardo DiCaprio",
9       "character": "Billy",
10      "main": true,
11      "urlCharacter": "http://www.imdb.com/character/ch0251381",
12      "urlProfile": "http://www.imdb.com/name/nm0000138"
13    },
14    {
15      "actorName": "Matt Damon",
16      "character": "Colin Sullivan",
17      "main": true,
18      "urlCharacter": "http://www.imdb.com/character/ch0002488",
19      "urlProfile": "http://www.imdb.com/name/nm0000354"
20    }
21  ]
2 }
```

# Document Store (Mapping)

RDBMS	MongoDB
Database	Database
Table, View	Collection
Row	JSON Document
Column	Field
Index	Index
Join	Embedded Document / Linking across Document
Foreign Key	Reference
Partition Key	Shard



# Sample Query – SQL vs. Mongo

SQL	Mongo
CREATE TABLE movies( movieId int NOT NULL AUTO_INCREMENT, name VARCHAR(30), rating VARCHAR(6), PRIMARY KEY (movieId ) )	db.movies.insert({ "id": 10, "name": "Titanic", "rating": "R" })
SELECT * FROM movies	db.movies.find()
UPDATE movies SET rating = "NR" WHERE movieId = 101	db.movies.update( {"id": 101 }, { \$set: { rating: "NR" } } )
DELETE FROM movies WHERE rating = "R"	db.movies.remove( { "rating": "R" } )



# Why MongoDB?

- ✧ “Queryable” documents
- ✧ No impedance mismatch between object and DB form
  - Ideal for web applications (fast retrieval)
- ✧ Quick and easy integration of new data variations
- ✧ Rich API support (multiple languages)





# Ruby On Rails & Mongo

## ✧ Ruby Driver

- <http://docs.mongodb.org/ecosystem/tutorial/ruby-driver-tutorial/>

## ✧ Mongoid

- <http://docs.mongodb.org/ecosystem/tutorial/ruby-mongoid-tutorial/>



# MongoDB Users

MetLife



Forbes



# MongoDB Core Topics With Ruby/Rails

- ✧ MongoDB Ruby Driver
- ✧ Aggregation Framework
- ✧ GridFS – breaking large files in to smaller chunks
- ✧ Geo Spatial - index and query geospatial data
- ✧ Mongoid



# Summary

- ✧ Open Source DB
- ✧ Automatic Scaling
- ✧ High Performance
- ✧ “Schema-less” and Document Oriented

## What's Next?

- ✧ MongoDB Installation



# In this lecture, we will discuss...

- ✧ Install MongoDB
- ✧ Configure MongoDB
- ✧ Start MongoDB - `mongod`
- ✧ Launch MongoDB shell - `mongo`

# MongoDB Installation Steps

<https://www.mongodb.org/downloads>

Download latest version

Supports all platforms

- ✧ Step 1: Download MongoDB (msi)
- ✧ Step 2: Mongo needs a default data folder (**very important step**)
  - Ex: /data/db or C:\data\db
  - **Note:** You can pick any folder but will need to provide the path while starting MongoDB



# Helpful Configuration

- ✧ Journaling in MongoDB – allocates 3GB upfront
  - Write-ahead logging to guarantee write operations
- ✧ For casual development, turn off (maybe?)
- ✧ Setting "nojournal=true" in mongod.conf will keep mongo from claiming this space for write-ahead journaling : `mongod --config /etc/mongod.conf`

Note: Do not turn off journaling in production system



# Starting MongoDB

- ✧ Open a CMD window and go to the `$mongo_install/bin` folder
- ✧ Step 1 : Start mongoDB
  - `mongod`
- ✧ Step 2 : Start mongo shell
  - `mongo` - [reference](#)
- ✧ *Note: If your db path is not the default, make sure to launch with this command - `mongod -dbpath /<path>`*
  - Directory **needs** to have **write permission**





# Summary

- ✧ MongoDB supports all OS
- ✧ MongoDB needs data folder
- ✧ Starting MongoDB - **mongod**
- ✧ Launching MongoDB shell – **mongo**

## What's Next?

- ✧ MongoDB Basics



# In this lecture, we will discuss...

- ✧ Importing sample data
- ✧ Basics of MongoDB shell
- ✧ MongoDB collections
- ✧ IRB shell and MongoDB
- ✧ Basic MongoDB commands in IRB

# MongoDB Basics

## ✧ Import dataset

- Download sample zips.json file from MongoDB
- Save the above file
- Run the import command as in

✧ `> mongoimport --db test --collection zips  
--drop --file zips.json`



# Database, Documents and Collections

- ✧ Mongo can create database **on the fly**
  - No need to create database beforehand
- ✧ Documents
  - Unit of **storing data** in a MongoDB database
  - JSON document
- ✧ Collection (similar to tables in DB)
  - Unit of **storing data** in a MongoDB database
  - Collection of documents

# Collection Types

## ✧ Capped Collection

- **Fixed-size** collections that support **high-throughput** operations
- Insert and retrieve documents based on **insertion order**
- Once a collection fills its allocated space, it **makes room** for new documents by **overwriting** the oldest documents in the collection
- `db.createCollection("log", { capped : true, size : 5242880, max : 5000 } )`



# Mongo Basics

## ✧ Start mongo shell

- `$ mongo`

## ✧ Switch to test database

```
> use test
```

## ✧ Test the data with a simple find command (note: we will cover this in more depth later)

```
> db.zips.findOne()
```

- The above command will return a single document from the zips collection.



# MongoDB Ruby Driver Setup

## ✧ mongo-ruby driver

- `gem update --system`
- `gem install mongo`
- `gem install bson_ext`

## ✧ Using gem

```
> require mongo
```



# MongoDB Basics (irb shell)

✧ Start irb shell

✧ Type the following commands:

```
> require 'mongo'
> Mongo::Logger.logger.level =
  ::Logger::INFO
> db =
Mongo::Client.new('mongodb://localhost:27017')
> db=db.use('test')
> db.database.name
> db.database.collection_names
> db[:zip].find.first
```





# Summary

- ✧ Basics of MongoDB
- ✧ Database, Document and Collection
- ✧ MongoDB Ruby Driver
  - `<irb>` shell

## What's Next?

- ✧ CRUD Operations



# Next Topic.....

## Lesson 2 – CRUD operations

